

# Pytorch

Ramex

June 2024

## 1 Tensors

A Tensor is like an n-dimensional array (multi-dimensional array) in numpy, but with additional features. A tensor can be created from a list or a numpy array. The tensor can be converted to a numpy array using the `numpy()` method. The tensor itself can be used for gpu computations.

### 1.1 Initializing a Tensor

A tensor can be initialized by a number of methods, for example:

- 1. Directly from data  
`data = [[1,2,3],[4,5,6]],`  
`data_ = torch.tensor(data)`
- 2. From a numpy array  
`data = np.array(data),`  
`data_x = torch.tensor(data)`
- 3. From another tensor (creates a copy of the tensor with ones only)  
`torch.ones_like(x_data)`
- 4. Random or constant values (this takes the shape as input)  
`x_ones = torch.ones_like(x_data)`

### 1.2 Attributes of a Tensor

A tensor has the following attributes:

- 1. `shape` : The shape of the tensor
- 2. `dtype` : The data type of the tensor
- 3. `device` : The device on which the tensor is stored
- 4. `size` : The number of elements in the tensor
- 5. `numel` : The number of elements in the tensor

- 6. `T` : The transposed tensor
- 7. `contiguous` : The contiguous tensor
- 8. `view` : The view of the tensor
- 9. `requires_grad` : The gradient required for the tensor
- 10. `grad` : The gradient of the tensor
- more ...  $\Rightarrow$  Pytorch Documentation

### 1.3 Indexing and Slicing

Indexing and slicing works the same as in numpy. For example:

- 1. `x[0]` : The first element of the tensor
- 2. `x[0,0]` : The first element of the first row
- 3. `x[0,:]` : The first row of the tensor
- 4. `x[:,0]` : The first column of the tensor
- 5. `x[0:2,0:2]` : The first two rows and columns of the tensor

### 1.4 Joining tensors

Tensors can be joined using the `torch.cat()` method. For example:

- 1. `torch.cat([x,y], dim=0)` : Concatenates the tensors along the rows
- 2. `torch.cat([x,y], dim=1)` : Concatenates the tensors along the columns
- 3. `torch.stack([x,y], dim=0)` : Stacks the tensors along the rows
- 4. `torch.stack([x,y], dim=1)` : Stacks the tensors along the columns

### 1.5 Single-element tensors

A single-element tensor is a tensor with one element. For example:

- 1. `x.item()` : Returns the value of the tensor as a python number
- 2. `x.tolist()` : Returns the value of the tensor as a python list
- 3. `x.numpy()` : Returns the value of the tensor as a numpy array
- 4. `x.to(device)` : Moves the tensor to the specified device

## 1.6 Tensor to NumPy array

A tensor can be converted to a numpy array using the `numpy()` method. Example:

- 1. `x.numpy()` : Converts the tensor to a numpy array
- 2. `x.cpu().numpy()` : Converts the tensor to a numpy array on the cpu
- 3. `x.cuda().numpy()` : Converts the tensor to a numpy array on the gpu

## 2 Datasets & DataLoaders

### 2.1 Loading a Dataset

Loading datasets in PyTorch involves using the `torch.utils.data` module, which provides utilities for efficiently loading and processing data. The key components include Dataset and DataLoader.

#### 2.1.1 Key Components

- 1. Dataset: An abstract class representing a dataset. You need to subclass this and implement two methods:
  - 1.1. `__len__`: Returns the size of the dataset.
  - 1.2. `__getitem__`: Supports indexing such that `dataset[i]` can be used to get the *i*th sample.
- 2. DataLoader: Combines a dataset and a sampler, and provides an iterable over the given dataset. It supports batching, shuffling, and parallel data loading.

#### 2.1.2 Example: Loading a Custom Dataset

Let's go through an example of creating a custom dataset and loading it using PyTorch.

#### 2.1.3 Step 1: Import Required Libraries

```
1 import torch
2 from torch.utils.data import Dataset, DataLoader
3 import pandas as pd
4 from sklearn.preprocessing import LabelEncoder
```

### 2.1.4 Step 2: Create a Custom Dataset

Assume we have a CSV file data.csv with the following structure:

```
1 text,label
2 "I love PyTorch", positive
3 "I hate bugs", negative
```

We will create a custom dataset to load this data.

```
1 class TextDataset(Dataset):
2     def __init__(self, csv_file):
3         self.data = pd.read_csv(csv_file)
4         self.texts = self.data['text'].values
5         self.labels = self.data['label'].values
6
7         # Encode labels as integers
8         self.label_encoder = LabelEncoder()
9         self.labels = self.label_encoder.fit_transform
10            (self.labels)
11
12     def __len__(self):
13         return len(self.texts)
14
15     def __getitem__(self, idx):
16         text = self.texts[idx]
17         label = self.labels[idx]
18         return text, label
```

### 2.1.5 Step 3: Instantiate the Dataset and DataLoader

```
1 # Create an instance of the dataset
2 dataset = TextDataset(csv_file='data.csv')
3
4 # Create a DataLoader for the dataset
5 dataloader = DataLoader(dataset, batch_size=2, shuffle
6     =True, num_workers=2)
```

### 2.1.6 Explanation of DataLoader Parameters

- dataset: The dataset from which to load the data.
- batch\_size: How many samples per batch to load.
- shuffle: Set to True to have the data reshuffled at every epoch.
- num\_workers: How many subprocesses to use for data loading. 0 means that the data will be loaded in the main process.

## 2.2 Step 4: Iterate Through the DataLoader

A dataset can be iterated over using a for loop. For example:

```
1     for batch in dataloader:
2         texts, labels = batch
3         print(texts)
4         print(labels)
```

## 2.3 Using Built-In Datasets

PyTorch also provides utilities for loading several standard datasets, such as MNIST, CIFAR-10, and ImageNet, through the torchvision package.

### 2.3.1 Example: Loading the MNIST Dataset

```
1 import torchvision.transforms as transforms
2 from torchvision.datasets import MNIST
3
4 # Define a transform to normalize the data
5 transform = transforms.Compose([
6     transforms.ToTensor(),
7     transforms.Normalize((0.1307,), (0.3081,))
8 ])
9
10 # Download and load the training dataset
11 train_dataset = MNIST(root='mnist_data', train=True,
12                        download=True, transform=transform)
13
14 # Create a DataLoader for the training dataset
15 train_loader = DataLoader(train_dataset, batch_size
16                           =64, shuffle=True, num_workers=2)
17
18 # Iterate through the DataLoader
19 for batch in train_loader:
20     images, labels = batch
21     print(images.shape)
22     print(labels)
23     break
```

## 3 Transforms

Transformers are a type of neural network architecture designed for handling sequential data, such as text. They have become a cornerstone of modern natural

language processing (NLP) due to their ability to capture long-range dependencies and parallelize training. Here's an explanation of Transformers and how to implement them using PyTorch.

### 3.1 Transformer Architecture

The Transformer model was introduced in the paper "Attention is All You Need" by Vaswani et al. in 2017. It consists of an encoder-decoder structure, where both the encoder and decoder are composed of a stack of identical layers.

### 3.2 Key Components

- 1. Multi-Head Self-Attention Mechanism: Allows the model to focus on different parts of the input sequence simultaneously.
- 2. Positional Encoding: Adds information about the position of words in the sequence.
- 3. Feed-Forward Neural Network: Applied to each position separately and identically.
- 4. Layer Normalization and Residual Connections: Improve training dynamics by normalizing intermediate layers and adding shortcuts to skip connections.

### 3.3 PyTorch Implementation

PyTorch provides a built-in module for the Transformer model through `torch.nn.Transformer`. Here's a step-by-step guide to implementing a basic Transformer model in PyTorch.

#### 3.3.1 Step 1: Import Required Libraries

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import math
```

#### Step 2: Positional Encoding

Positional encoding helps the model to understand the order of the sequence.

```
1 class PositionalEncoding(nn.Module):
2     def __init__(self, d_model, max_len=5000):
3         super(PositionalEncoding, self).__init__()
```

```

4         self.pe = torch.zeros(max_len, d_model).
            unsqueeze(0)
5         position = torch.arange(0, max_len).unsqueeze
            (1)
6         div_term = torch.exp(torch.arange(0, d_model,
            2) * -(math.log(10000.0) / d_model))
7         self.pe[:, :, 0::2] = torch.sin(position *
            div_term)
8         self.pe[:, :, 1::2] = torch.cos(position *
            div_term)
9
10    def forward(self, x):
11        x = x + self.pe[:, :x.size(1)]
12        return x

```

### Step 3: Transformer Model

```

1    class TransformerModel(nn.Module):
2    def __init__(self, input_dim, model_dim,
        output_dim, nhead, num_encoder_layers,
        num_decoder_layers, dim_feedforward,
        max_seq_length):
3        super(TransformerModel, self).__init__()
4        self.model_type = 'Transformer'
5        self.pos_encoder = PositionalEncoding(
            model_dim, max_seq_length)
6        self.encoder = nn.Embedding(input_dim,
            model_dim)
7        self.decoder = nn.Embedding(output_dim,
            model_dim)
8        self.transformer = nn.Transformer(d_model=
            model_dim, nhead=nhead,
9            num_encoder_layers=num_encoder_layers,
10           num_decoder_layers=num_decoder_layers,
            dim_feedforward=dim_feedforward)
11        self.fc_out = nn.Linear(model_dim, output_dim)
12        self.model_dim = model_dim
13
14    def forward(self, src, tgt, src_mask=None,
        tgt_mask=None):
15        src = self.encoder(src) * math.sqrt(self.
            model_dim)
16        tgt = self.decoder(tgt) * math.sqrt(self.
            model_dim)
17        src = self.pos_encoder(src)

```

```

18         tgt = self.pos_encoder(tgt)
19         output = self.transformer(src, tgt, src_mask,
20                                   tgt_mask)
21         output = self.fc_out(output)
22         return output

```

#### Step 4: Example Usage

```

1     # Define the model parameters
2     input_dim = 1000 # Size of the input vocabulary
3     model_dim = 512  # Dimension of the model
4     output_dim = 1000 # Size of the output vocabulary
5     nhead = 8        # Number of attention heads
6     num_encoder_layers = 6
7     num_decoder_layers = 6
8     dim_feedforward = 2048
9     max_seq_length = 100
10
11     # Instantiate the model
12     model = TransformerModel(input_dim, model_dim,
13                               output_dim, nhead, num_encoder_layers,
14                               num_decoder_layers, dim_feedforward, max_seq_length
15                               )
16
17     # Define input and output sequences (batch_size,
18     # sequence_length)
19     src = torch.randint(0, input_dim, (10, 32)) # Example
20     # source sequence
21     tgt = torch.randint(0, output_dim, (10, 32)) # Example
22     # target sequence
23
24     # Forward pass
25     output = model(src, tgt)
26
27     print(output.shape) # Output shape will be (
28     # sequence_length, batch_size, output_dim)

```